



2019



# Data Science and AI

Module 2

Part 1:

---

SQL and Databases

---



# Agenda: Module 2 Part 1

- Introduction to Databases
- The relational database paradigm
- Basic SQL
- RDBMS
- Advanced SQL
- SQL in Python
- NoSQL Databases



# Introduction to Databases

- Databases: definition, usage, features, applications
- Database Elements
- Database Principles
- The relational database paradigm
- SQL
- RDBMS

# Introduction to Databases


- What is a database?
  - a computer system that manages **storage** and **querying** of data
- How is a database used?
  - data **insertion** & **retrieval** are (typically) performed using a **query language**
    - a compact programming syntax
    - basic operators for data transformation
- What are the essential features of a database?
  - practical design for organising data
  - **efficient** methods to retrieve specific information
  - indexing, **performance optimisation**
  - **reliability, security, backup & replication**

# Good News: Demand for Data Science!

**Demand for deep analytical talent in the United States could be 50 to 60 percent greater than its projected supply by 2018**

Supply and demand of deep analytical talent by 2018

Thousand people



<sup>1</sup> Other supply drivers include attrition (-), immigration (+), and reemploying previously unemployed deep analytical talent (+).

SOURCE: US Bureau of Labor Statistics; US Census; Dun & Bradstreet; company interviews; McKinsey Global Institute analysis

# Why Use a Database?

- The **standard** solution for data storage
- Much more **robust** than text, CSV or JSON files
- Most analyses involve pulling data to and from a resource; in most settings, this means using a **database**
- **Many types and variants** to serve different use cases
- Rules on structure make writing and retrieving data more **reliable and efficient**
- Provide a **central source of “truth”**

# Database Application Areas - Examples

- Operations

- transaction systems
- data capture
- inventory management

- Master Data

- products
- customers
- suppliers

- Data Warehouse

- Reporting
- Analytics
- Data Science



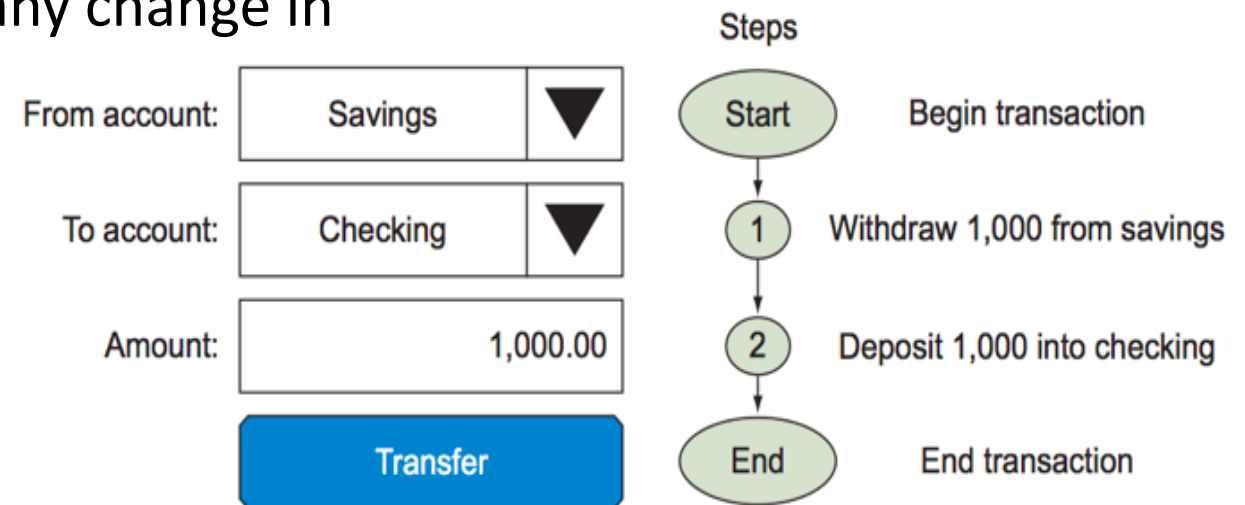
# Database Elements

- tables
  - data storage by **columns** (attributes) and **rows** (records)
- keys
  - for matching **indexed** attributes across tables
- queries, views
  - for retrieving, subsetting, aggregating, joining data
- functions, procedures
  - **reusable** code units
- types
  - reusable data structures
- triggers
  - procedures that run automatically when a specific **event** occurs
- jobs
  - batches of procedures that run on a **schedule**

# Database Principles: Transactional Integrity

*def:* Transaction

- a unit of work performed against a database
- this term generally represents any change in database
- involves **multiple steps**



# Database Principles: ACID

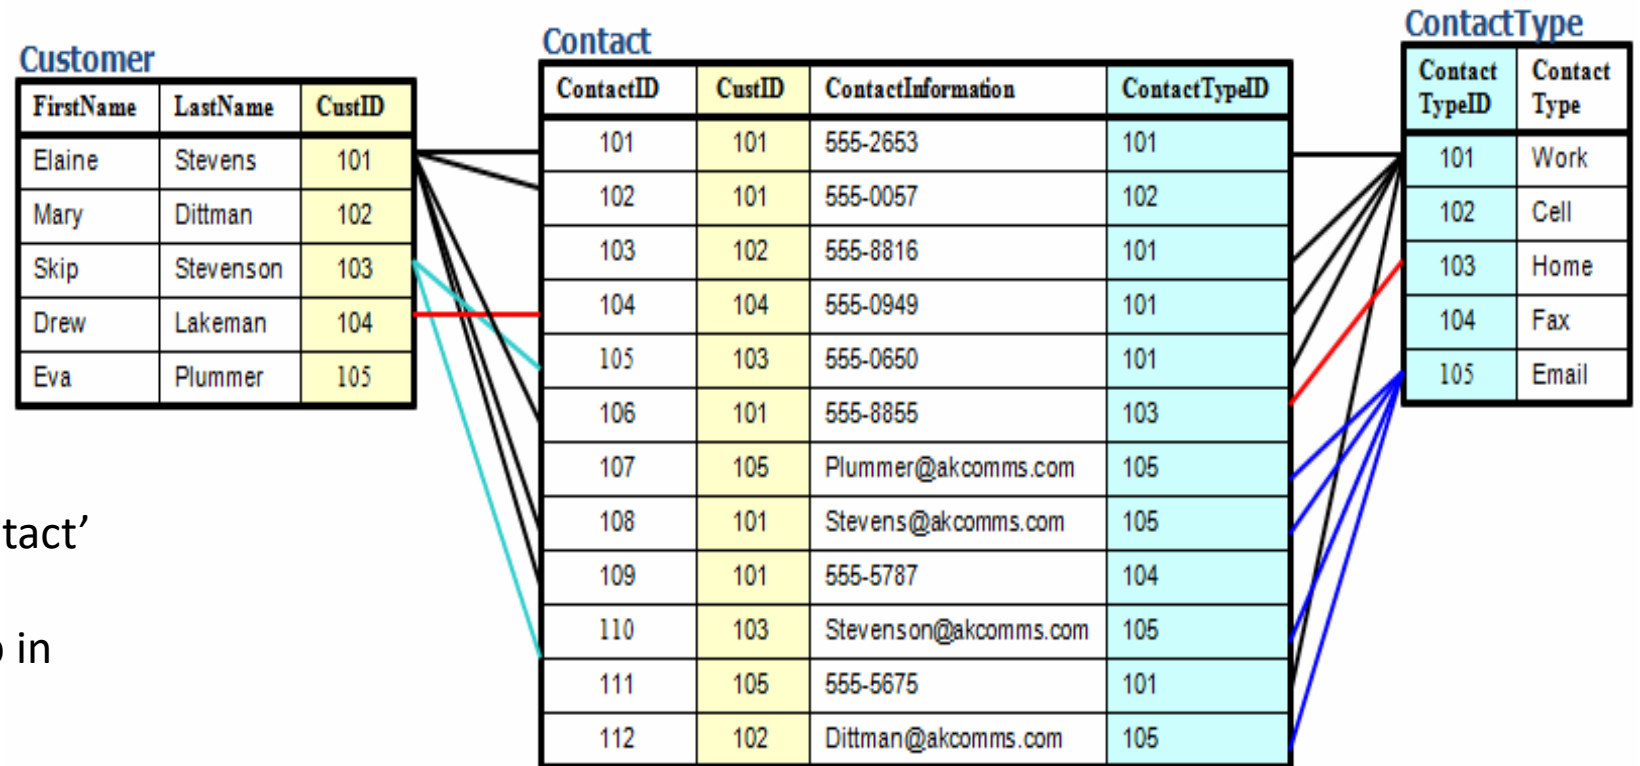
*def:* ACID is a set of properties that guarantee that database transactions are processed reliably



- **Atomicity:** if one part of the transaction fails, the entire transaction fails; the database state is left unchanged (**'all or nothing'**)
- **Consistency:** ensures that any transaction will bring the database from one **valid state** to another
- **Isolation:** ensures that the **concurrent** execution of transactions results in a system state that would be obtained if transactions were executed serially (one after the other)
- **Durability:** ensures that once a transaction has been committed it will be unaffected by **power loss, system crashes, or errors**

# The relational Database Paradigm

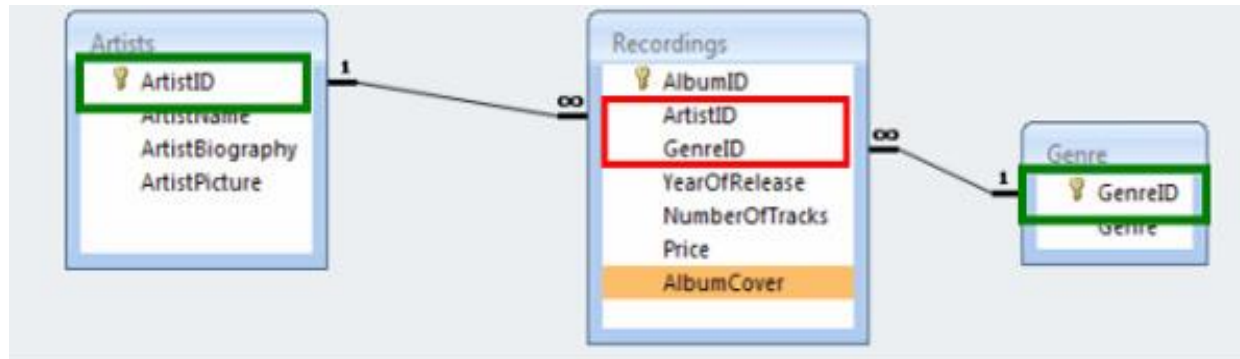
- Each table in a database is devoted to one **domain**
- Keys are used to **connect** tables in a logical manner



- customer info in 'Customer' table
- contacts info in 'Contact' table
- contact type lookup in 'ContactType' table

# Relational Databases

- Queries can pull together information from multiple tables by matching **foreign keys** to primary keys



= primary key



= foreign key

- 'ArtistID' is:
  - a foreign key in the 'Recordings' table
  - the primary key in the 'Artists' table
- 'GenreID' is:
  - a foreign key in the 'Recordings' table
  - the primary key in the 'Genre' table

# RDBMS

- RDB = ?
  - relational database
- RDBMS = ?
  - relational database management system



ORACLE®

PostgreSQL



SYBASE®



# RDBMS

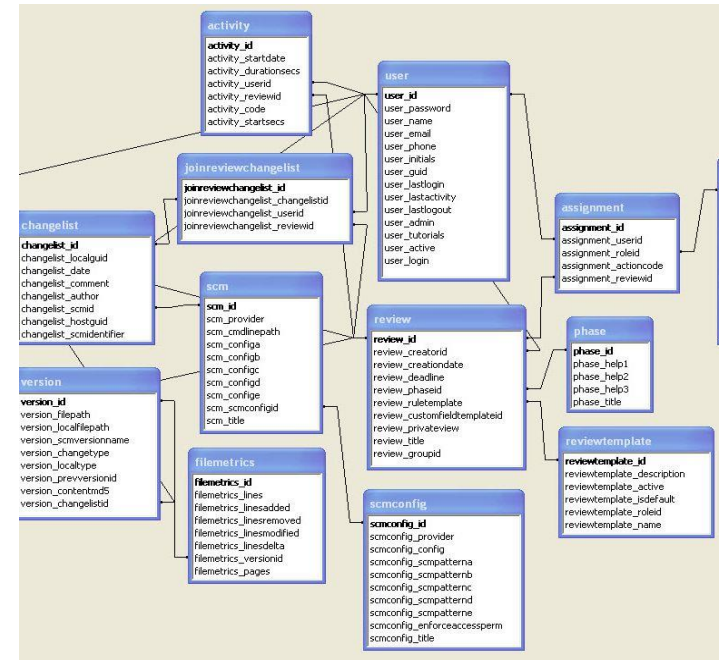
The RDBMS is the system, and SQL is the language used to interact with the system.

In principle you could have an RDBMS that uses some other language for access, and in principle you could use SQL to interact with some other kind of database system, though in practice the two are closely coupled.

SQLite, MySQL, MariaDB, Postgres, et al are all RDBMSes, and the language you'll use to interact with all of them is SQL.

# Database Schema

- table-level
  - columns and primary key
- database-level
  - overall design
    - data model (tables)
    - keys (PK, FK)
  - integrity constraints
- rationale
  - removes tight coupling of database objects and owners
  - improves security administration of database objects






# Relation Variables

By using a Data Definition Language (DDL), it is able to define base relation variables. In SQL, **CREATE TABLE** syntax is used to define base tables. The following is an example.

```
CREATE TABLE List_of_people (
  ID INTEGER,
  Name CHAR(40),
  Address CHAR(200),
  PRIMARY KEY (ID)
)
```

The Data Definition Language (DDL) is also used to define derived relation variables. In SQL, **CREATE VIEW** syntax is used to define a **derived relation variable**. The following is an example.

```
CREATE VIEW List_of_Okinawa_people AS (
  SELECT ID, Name, Address
  FROM List_of_people
  WHERE Address LIKE '%, Okinawa'
)
```



# SQL (Structured Query Language)

- essentially a **declarative** language
  - data are typically retrieved as **subsets** of the rows and columns in one or more tables: 'rowsets'
  - the user encodes the kind of rowset to be returned
  - the database engine produces a **plan** for how to execute it
    - **optimised** if possible
    - may accept **hints** from user (e.g. indexes to use)
- 3 functional divisions:
  - data definition language (**DDL**)
    - schema creation and modification
  - data manipulation language (**DML**)
    - insert, update, delete
  - data control language (**DCL**)
    - access, security

## SQL – cont'd

simple rowset retrieval:

```
SELECT columns FROM table
```

conditional rowset retrieval:

```
SELECT columns FROM table WHERE condition
```

*example (conditional rowset with ordering):*

```
SELECT TOP 10 amount FROM sales WHERE amount > 99.99 ORDER BY amount
```

# SQL – cont'd

## Syntax

- identifiers with embedded **spaces or punctuation** require implementation-specific delimiters
  - e.g. SQL Server allows “my table” or [my table]  
(delimiters are optional, otherwise)
- **namespace hierarchies** are usually delimited with periods (full stops)
  - e.g. (SQL Server): myschema.mytable.mycolumn

# SQL – cont'd

## Syntax

- SQL is **case-insensitive**
  - best to adopt a style, for readability  
(e.g. uppercase for SQL keywords, lowercase for identifiers)
- Microsoft and others extend the ANSI standard
  - extensions add **power and convenience** to programming
  - many extensions do not readily **port** between versions of SQL

# Indices (*aka* 'Indexes')

- An index is a column of unique numbers (one per row) used to speed up query performance
  - reduces number of database data pages that have to be **scanned**
  - can have **> 1** index per table
  - usually based on single columns or **tuples** of columns
- a **clustered index** (e.g. SQL Server) determines physical order of data in a table
  - can only have 1 clustered index per table
- Building / rebuilding an index is an **expensive operation**
  - for inserting a large number of rows, it is usually best to **drop the index**, do the insert, then rebuild the index

# Primary and Foreign Keys

- *primary key* uniquely refers to each row in a table
  - can have only one per table
  - usually an integer
- *foreign key* refers to a row in a different table
  - is a primary key in the other table
  - can have any number per table
- Nb. the use of foreign keys reduces storage and makes database maintenance much easier

# Some samples

*w3schools:*

- *<https://www.w3schools.com/sql/>*



# Some of The Most Important SQL Commands

**SELECT** - *extracts data from a database*

**UPDATE** - *updates data in a database*

**DELETE** - *deletes data from a database*

**INSERT INTO** - *inserts new data into a database*

**CREATE DATABASE** - *creates a new database*

**ALTER DATABASE** - *modifies a database*

**CREATE TABLE** - *creates a new table*

**ALTER TABLE** - *modifies a table*

**DROP TABLE** - *deletes a table*

**CREATE INDEX** - *creates an index (search key)*

**DROP INDEX** - *deletes an index*

# SELECT & SELECT DISTINCT

*The **SELECT** statement is used to select data from a database. The data returned is stored in a result table, called the result-set.*

**SELECT** Syntax

**SELECT** column1, column2, ...  
**FROM** table\_name;

**SELECT** CustomerName, City **FROM** Customers;

*The following SQL statement lists the number of different (distinct) customer countries:*

*The **SELECT DISTINCT** statement is used to return only distinct (different) values.*

*Inside a table, a column often contains many duplicate values; and sometimes you only want to list the different (distinct) values.*

**SELECT DISTINCT** Syntax

**SELECT DISTINCT** column1, column2, ...  
**FROM** table\_name;

**SELECT DISTINCT** Country **FROM** Customers;  
**SELECT COUNT(DISTINCT** Country **FROM** Customers;

**SELECT** Count(\*) AS DistinctCountries  
**FROM** (SELECT DISTINCT Country **FROM** Customers);

# WHERE

## The SQL WHERE Clause

*The WHERE clause is used to filter records.*

*The WHERE clause is used to extract only those records that fulfill a specified condition.*

### WHERE Syntax

```
SELECT column1, column2, ...
FROM table_name
WHERE condition;
```

### Text Fields vs. Numeric Fields

```
SELECT * FROM Customers
WHERE Country='Mexico';
```

```
SELECT * FROM Customers
WHERE CustomerID=1;
```

## Operators in The WHERE Clause

*The following operators can be used in the WHERE clause:*

Operator	Description	Example
=	Equal	
>	Greater than	
<	Less than	
>=	Greater than or equal	
<=	Less than or equal	
<>	Not equal. Note: In some versions of SQL this operator may be written as !=	
BETWEEN	Between a certain range	
LIKE	Search for a pattern	
IN	To specify multiple possible values for a column	

# INSERT

*The **INSERT INTO** statement is used to insert new records in a table.*

## **INSERT INTO Syntax**

*It is possible to write the **INSERT INTO** statement in two ways.*

*The first way specifies both the column names and the values to be inserted:*

```
INSERT INTO table_name (column1, column2,  
column3, ...)  
VALUES (value1, value2, value3, ...);
```

```
INSERT INTO table_name  
VALUES (value1, value2, value3, ...);
```

## **INSERT INTO Example**

*The following SQL statement inserts a new record in the "Customers" table:*

### **Example**

```
INSERT INTO Customers (CustomerName,  
ContactName, Address, City, PostalCode, Country)  
VALUES ('Cardinal', 'Tom B. Erichsen', 'Skagen 21',  
'Stavanger', '4006', 'Norway');
```

### **Insert Data Only in Specified Columns**

```
INSERT INTO Customers (CustomerName, City,  
Country)  
VALUES ('Cardinal', 'Stavanger', 'Norway');
```

# UPDATE

*The **UPDATE** statement is used to modify the existing records in a table.*

**UPDATE Syntax**

**UPDATE** *table\_name*

**SET** *column1 = value1, column2 = value2, ...*

**WHERE** *condition;*

**UPDATE Customers**

**SET** *ContactName = 'Alfred Schmidt', City= 'Frankfurt'*

**WHERE** *CustomerID = 1;*

**UPDATE Multiple Records**

*It is the **WHERE** clause that determines how many records that will be updated.*

**UPDATE Customers**

**SET** *ContactName='Juan'*

**WHERE** *Country='Mexico';*

**Update Warning!**

*Be careful when updating records. If you omit the **WHERE** clause, ALL records will be updated!*

**Example**

**UPDATE Customers**

**SET** *ContactName='Juan';*

# DELETE

*The **DELETE** statement is used to delete existing records in a table.*

**DELETE Syntax**

**DELETE FROM** *table\_name* **WHERE** *condition*;

*Example*

**DELETE FROM** *Customers* **WHERE** *CustomerName='Alfreds Futterkiste'*;

*The following SQL statement deletes all rows in the "Customers" table, without deleting the table:*

*Example*

**DELETE FROM** *Customers*;

# SELECT TOP

## *The SQL **SELECT TOP** Clause*

*The **SELECT TOP** clause is used to specify the number of records to return.*

*The **SELECT TOP** clause is useful on large tables with thousands of records. Returning a large number of records can impact on performance.*

*Note: Not all database systems support the **SELECT TOP** clause. MySQL supports the **LIMIT** clause to select a limited number of records, while Oracle uses **ROWNUM**.*

## *SQL Server / MS Access Syntax:*

**SELECT TOP** number|percent column\_name(s)  
**FROM** table\_name  
**WHERE** condition;

## *MySQL Syntax:*

**SELECT** column\_name(s)  
**FROM** table\_name  
**WHERE** condition  
**LIMIT** number;

## *Oracle Syntax:*

**SELECT** column\_name(s)  
**FROM** table\_name  
**WHERE ROWNUM** <= number;

## *SQL TOP, LIMIT and ROWNUM Examples*

**SELECT \* FROM** Customers  
**LIMIT 3;**  
**SELECT \* FROM** Customers  
**WHERE ROWNUM** <= 3;  
**SELECT TOP 50 PERCENT \* FROM** Customers;

# CREATE DATABASE TABLE

*The **CREATE DATABASE** statement is used to create a new SQL database.*

*Syntax*

**CREATE DATABASE** *databasename*;

**CREATE DATABASE** *testDB*;

*The **CREATE TABLE** statement is used to create a new table in a database.*

*Syntax*

**CREATE TABLE** *table\_name* (  
    *column1 datatype,*  
    *column2 datatype,*  
    *column3 datatype,*  
    ....  
);

**CREATE TABLE** *Persons* (  
    *PersonID int,*  
    *LastName varchar(255),*  
    *FirstName varchar(255),*  
    *Address varchar(255),*  
    *City varchar(255)*  
);

*The following SQL creates a new table called "TestTables" (which is a copy of the "Customers" table):*

*Example*

**CREATE TABLE** *TestTable* **AS**  
**SELECT** *customername, contactname*  
**FROM** *customers*;



# ALTER TABLE

*The ALTER TABLE statement is used to add, delete, or modify columns in an existing table. The ALTER TABLE statement is also used to add and drop various constraints on an existing table.*

## **ALTER TABLE - ADD Column**

*To add a column in a table, use the following syntax:*

```
ALTER TABLE table_name  
ADD column_name datatype;
```

```
ALTER TABLE Customers  
ADD Email varchar(255);
```

```
ALTER TABLE table_name  
DROP COLUMN column_name;
```

```
ALTER TABLE Customers  
DROP COLUMN Email;
```

*To change the data type of a column in a table, use the following syntax:*

*SQL Server / MS Access:*

```
ALTER TABLE table_name  
ALTER COLUMN column_name datatype;
```

*My SQL / Oracle (prior version 10G):*

```
ALTER TABLE table_name  
MODIFY COLUMN column_name datatype;
```

*Oracle 10G and later:*

```
ALTER TABLE table_name  
MODIFY column_name datatype;
```

```
ALTER TABLE Persons  
ADD DateOfBirth date;
```

# PRIMARY KEY

The **PRIMARY KEY** constraint uniquely identifies each record in a table.

Primary keys must contain **UNIQUE** values, and cannot contain NULL values.

A table can have only ONE primary key; and in the table, this primary key can consist of single or multiple columns (fields).

**SQL PRIMARY KEY on CREATE TABLE**

The following SQL creates a PRIMARY KEY on the "ID" column when the "Persons" table is created:

MySQL:

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int,  
    PRIMARY KEY (ID)  
);
```

SQL Server / Oracle / MS Access:

```
CREATE TABLE Persons (  
    ID int NOT NULL PRIMARY KEY,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int  
);
```

To allow naming of a PRIMARY KEY constraint, and for defining a PRIMARY KEY constraint on multiple columns, use the following SQL syntax:

MySQL / SQL Server / Oracle / MS Access:

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int,  
    CONSTRAINT PK_Person PRIMARY KEY  
    (ID,LastName)  
);
```

# CREATE INDEX

*The **CREATE INDEX** statement is used to create indexes in tables.*

*Indexes are used to retrieve data from the database very fast. The users cannot see the indexes, they are just used to speed up searches/queries.*

*Note: Updating a table with indexes takes more time than updating a table without (because the indexes also need an update). So, only create indexes on columns that will be frequently searched against.*

## **CREATE INDEX** Syntax

*Creates an index on a table. Duplicate values are allowed:*

**CREATE INDEX** *index\_name*

**ON** *table\_name* (*column1*, *column2*, ...);

## **CREATE UNIQUE INDEX** Syntax

*Creates a unique index on a table. Duplicate values are not allowed:*

**CREATE UNIQUE INDEX** *index\_name*

**ON** *table\_name* (*column1*, *column2*, ...);

# CREATE INDEX - cont'd

**CREATE INDEX** *idx\_lastname*

**ON** *Persons (LastName);*

*If you want to create an index on a combination of columns, you can list the column names within the parentheses, separated by commas:*

**CREATE INDEX** *idx\_pname*

**ON** *Persons (LastName, FirstName);*

**DROP INDEX** *Statement*

*The **DROP INDEX** statement is used to delete an index in a table.*

*MS Access:*

**DROP INDEX** *index\_name* **ON** *table\_name;*

*SQL Server:*

**DROP INDEX** *table\_name.index\_name;*

*DB2/Oracle:*

**DROP INDEX** *index\_name;*

*MySQL:*

**ALTER TABLE** *table\_name*

**DROP INDEX** *index\_name;*

# DROP DATABASE TABLE

*The **DROP DATABASE** statement is used to drop an existing SQL database.*

## **DROP DATABASE** Example

*The following SQL statement drops the existing database "testDB":*

*Example*

**DROP DATABASE** testDB;

*The **DROP TABLE** statement is used to drop an existing table in a database.*

## **SQL DROP TABLE** Example

*The following SQL statement drops the existing table "Shippers":*

*Example*


**DROP TABLE** Shippers;

# Joins


- used (in queries) for
  - **combining columns** from different tables
  - **filtering columns** from one table using criteria in a different table
- joins match a foreign key in one table to a primary key in another table
- Nb. simple joins are generally quite fast, but **compound joins** (involving many tables) can be very slow
- a database join is analogous to a set operation in mathematics

# Joins


## SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```





```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

© C.L. Moffatt, 2008

# Joins – cont'd

- Compound inner join example



Return artist name and album cover for every album of the 'rock' genre:

```
SELECT Artists.ArtistName,  
       Recordings.AlbumCover  
FROM Artists  
INNER JOIN Recordings  
ON Artists.ArtistID = Recordings.ArtistID  
INNER JOIN Genre  
ON Recordings.GenreID = Genre.GenreID  
WHERE Genre.Genre = 'rock'
```



# Joins – cont'd

## Left Join

- all rows from 1<sup>st</sup> table, plus matching rows from 2<sup>nd</sup> table:

```
SELECT cars.car, trucks.truck FROM cars
LEFT JOIN trucks ON cars.maker = trucks.maker
```

- unmatched rows will have NULL in place of truck



## Right Join

- all rows from 2<sup>nd</sup> table, plus matching rows from 1<sup>st</sup> table:

```
SELECT cars.car, trucks.truck FROM cars
RIGHT JOIN trucks ON cars.maker = trucks.maker
```


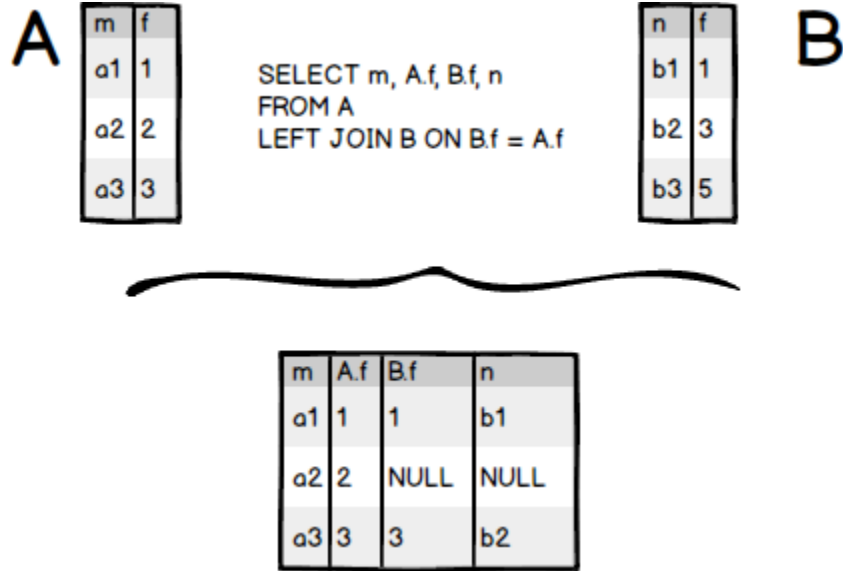
- unmatched rows will have NULL in place of car

Table 1	Table 2
aaa	aaa
bbb	xxx
ccc	ccc
ddd	yyy
eee	eee
fff	fff



# Joins – cont'd

## Left Join



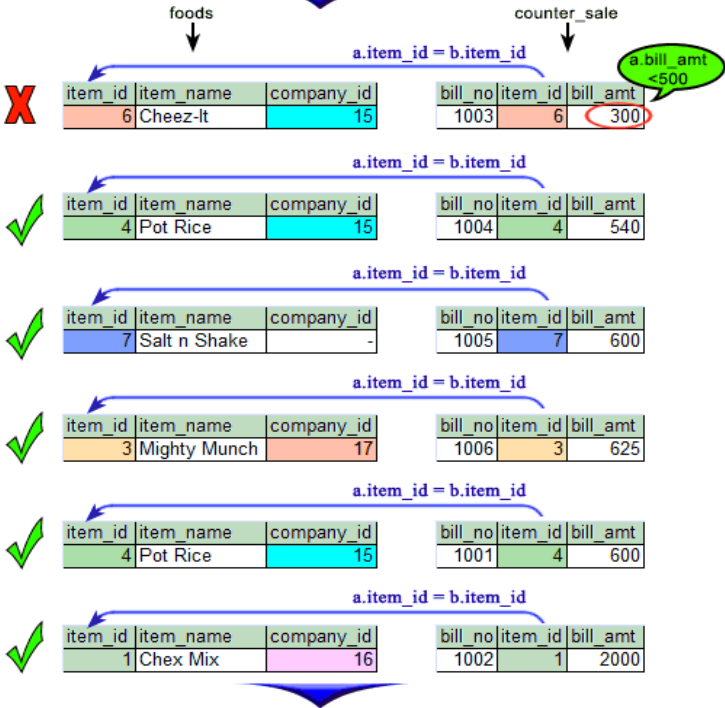
# Joins – cont’d

## sql left join on multiple columns

This SQL statement will **first join all rows** from the **counter\_sale** table and **only** those rows from the **foods** table where the joined fields are **equal** and if the **ON clause matches** no records in the **foods** table, the join will still return rows, but the **NULL** in each column of right table, therefore eliminates those rows which bill amount is less than or equal to 500.

```
SELECT a.bill_no, b.item_name, a.bill_amt
FROM counter_sale a
LEFT JOIN foods b
ON a.item_id=b.item_id
WHERE a.bill_amt>500;
```

foods (b)				counter_sale (a)					
item_id	item_name	item_unit	company_id	bill_no	item_id	sl_qty	sl_rate	bill_amt	
1	Chex Mix	Pcs	16	1003	6	15	20	300	
6	Cheez-It	Pcs	15	1004	4	18	30	540	
2	BN Biscuit	Pcs	15	1005	7	10	60	600	
3	Mighty Munch	Pcs	17	1006	3	25	25	625	
4	Pot Rice	Pcs	15	1001	4	20	30	600	
5	Jaffa Cakes	Pcs	18	1002	1	40	50	2000	
7	Salt n Shake	Pcs	-						



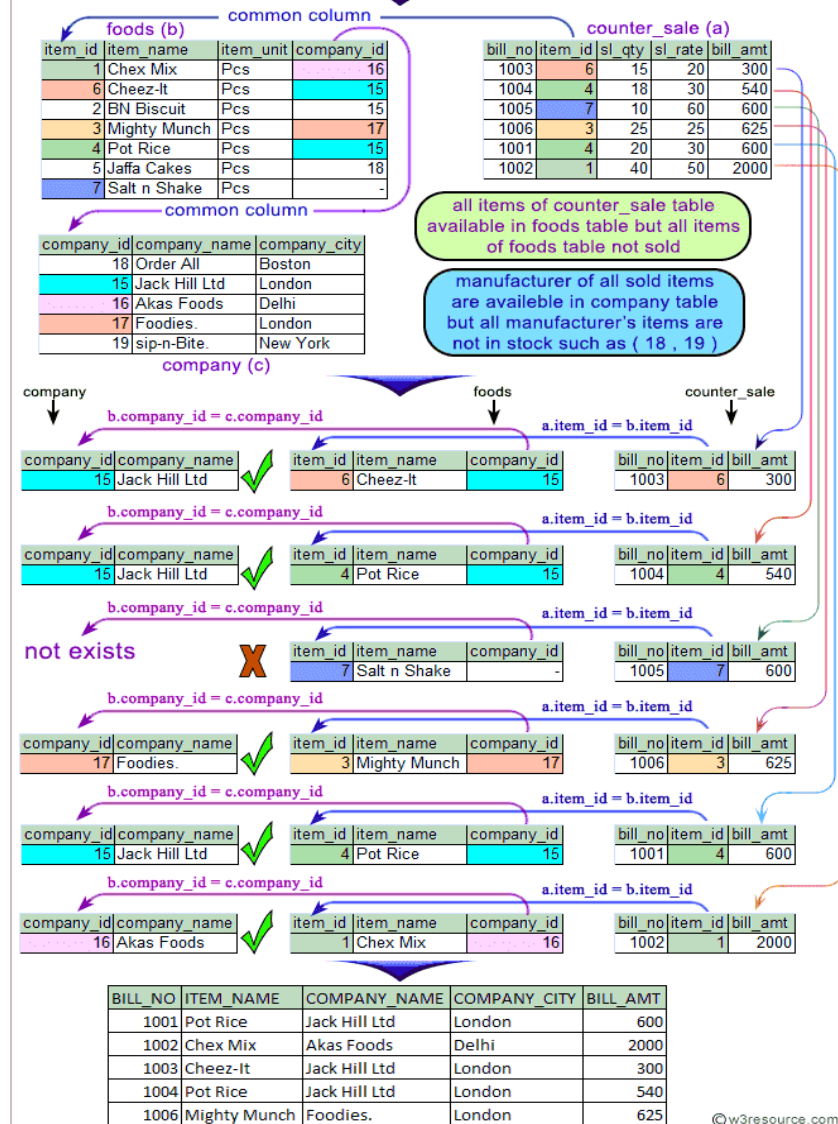
BILL_NO	ITEM_NAME	BILL_AMT
1002	Chex Mix	2000
1006	Mighty Munch	625
1001	Pot Rice	600
1004	Pot Rice	540
1005	Salt n Shake	600

# Joins – cont'd

## Left Join on multiple table

To filtered out those bill number, item name, company name and city and the bill amount for each bill, which items are available in foods table, and their manufacturer must have enlisted to supply that item, and no NULL value for manufacturer are not allowed

```
SELECT a.bill_no, b.item_name,
       c.company_name, c.company_city, a.bill_amt
FROM counter_sale a
LEFT JOIN foods b ON a.item_id=b.item_id
LEFT JOIN company c on b.company_id=c.company_id
WHERE c.company_name IS NOT NULL
ORDER BY a.bill_no;
```




# Joins – cont'd

## Outer Join

- all rows from 1<sup>st</sup> table, matched with all rows from 2<sup>nd</sup> table:

```
SELECT cars.car, trucks.truck FROM cars  
OUTER JOIN trucks ON cars.maker = trucks.maker
```

- returns every possible pairing of car and truck
- uses:
  - creating contingency tables
  - creating dummy data for testing
  - other?



# Joins – cont'd

**Unfortunately, SQLite** does not support the **RIGHT JOIN** clause and also the **FULL OUTER JOIN** clause. However, you can easily emulate the **FULL OUTER JOIN** by using the **LEFT JOIN** clause.

```
-- create and insert data into the dogs table
CREATE TABLE dogs (
  type    TEXT,
  color TEXT
);
INSERT INTO dogs(type, color)
VALUES('Hunting','Black'), ('Guard','Brown');

-- create and insert data into the cats table
CREATE TABLE cats (
  type    TEXT,
  color TEXT
);
```

```
INSERT INTO cats(type,color)
VALUES('Indoor','White'),
      ('Outdoor','Black');
```

```
SELECT d.type,
       d.color,
       c.type,
       c.color
FROM dogs d
LEFT JOIN cats c USING(color)
UNION ALL
SELECT d.type,
       d.color,
       c.type,
       c.color
FROM cats c
LEFT JOIN dogs d ON d.color=c.color
WHERE d.color IS NULL;
```

# Writing code using DB-API - Update Operation

```
import MySQLdb
# Open database connection
db = MySQLdb.connect("localhost","testuser","test123","TESTDB" )
# prepare a cursor object using cursor() method
cursor = db.cursor()
# Drop table if it already exist using execute() method.
cursor.execute("DROP TABLE IF EXISTS EMPLOYEE")

# Create table as per requirement
sql = """CREATE TABLE EMPLOYEE (
    FIRST_NAME  CHAR(20) NOT NULL,
    LAST_NAME   CHAR(20),
    AGE INT,
    SEX CHAR(1),
    INCOME FLOAT )"""

cursor.execute(sql)
# disconnect from server
db.close()
```

# Connect to Database and Create Table

If the database does not exist, then it will be created and finally a database object will be returned.

```
#!/usr/bin/python
import sqlite3
conn = sqlite3.connect('test.db')
print "Opened database successfully";

#Create table
conn.execute("""CREATE TABLE COMPANY
      (ID INT PRIMARY KEY     NOT NULL,
      NAME           TEXT     NOT NULL,
      AGE            INT       NOT NULL,
      ADDRESS        CHAR(50),
      SALARY         REAL);""")
print "Table created successfully";
conn.close()
```



# Writing code using sqlite3-API - Insert Operation

```
#!/usr/bin/python
import sqlite3
conn = sqlite3.connect('test.db')
print "Opened database successfully";
conn.execute("INSERT INTO COMPANY (ID,NAME,AGE,ADDRESS,SALARY) \
VALUES (1, 'Paul', 32, 'California', 20000.00 )");
conn.execute("INSERT INTO COMPANY (ID,NAME,AGE,ADDRESS,SALARY) \
VALUES (2, 'Allen', 25, 'Texas', 15000.00 )");
conn.execute("INSERT INTO COMPANY (ID,NAME,AGE,ADDRESS,SALARY) \
VALUES (3, 'Teddy', 23, 'Norway', 20000.00 )");
conn.execute("INSERT INTO COMPANY (ID,NAME,AGE,ADDRESS,SALARY) \
VALUES (4, 'Mark', 25, 'Rich-Mond ', 65000.00 )");
conn.commit()
print "Records created successfully";
conn.close()
```

# Writing code using sqlite3-API - Select Operation

```
#!/usr/bin/python
import sqlite3
conn = sqlite3.connect('test.db')
print "Opened database successfully";
cursor = conn.execute("SELECT id, name, address, salary from COMPANY")
for row in cursor:
    print "ID = ", row[0]
    print "NAME = ", row[1]
    print "ADDRESS = ", row[2]
    print "SALARY = ", row[3], "\n"
print "Operation done successfully";
conn.close()
```

# Writing code using sqlite3-API - update Operation

```
#!/usr/bin/python
import sqlite3
conn = sqlite3.connect('test.db')
print "Opened database successfully";
conn.execute("UPDATE COMPANY set SALARY = 25000.00 where ID = 1")
conn.commit
print "Total number of rows updated :", conn.total_changes
cursor = conn.execute("SELECT id, name, address, salary from COMPANY")
for row in cursor:
    print "ID = ", row[0]
    print "NAME = ", row[1]
    print "ADDRESS = ", row[2]
    print "SALARY = ", row[3], "\n"
print "Operation done successfully";
conn.close()
```

# Writing code using sqlite3-API - DELETE Operation

```
#!/usr/bin/python
import sqlite3
conn = sqlite3.connect('test.db')
print "Opened database successfully";
conn.execute("DELETE from COMPANY where ID = 2;")
conn.commit()
print "Total number of rows deleted :", conn.total_changes
cursor = conn.execute("SELECT id, name, address, salary from COMPANY")
for row in cursor:
    print "ID = ", row[0]
    print "NAME = ", row[1]
    print "ADDRESS = ", row[2]
    print "SALARY = ", row[3], "\n"
print "Operation done successfully";
conn.close()
```

[https://www.tutorialspoint.com/sqlite/sqlite\\_quick\\_guide.htm](https://www.tutorialspoint.com/sqlite/sqlite_quick_guide.htm)  
<https://gist.github.com/naiquevin/1746005>

# Connecting to a database using ibm\_db API

The `ibm_db` API provides a variety of useful Python functions for accessing and manipulating data in an IBM data server database, including functions for connecting to a database, preparing and issuing SQL statements, fetching rows from result sets, calling stored procedures, committing and rolling back transactions, handling errors and retrieving metadata.

The `ibm_db` API uses the IBM Data Server Driver for ODBC, and CLI APIs to connect to IBM, DB2 and Informix.

<http://cican17.com/accessing-database-using-python/>

# SQLite Aggregate Functions -Average

The AVG function is an aggregate function that calculates the average value of all non-NULL values within a group.

```
CREATE TABLE avg_tests (val);
INSERT INTO avg_tests (val)
VALUES (1), (2), (10.1), (20.5),
('8'), ('B'), (NULL), (x'0010'), (x'0011');
```

```
SELECT rowid,
       val
FROM avg_tests;
```

```
SELECT avg(val)
FROM avg_tests
WHERE rowid < 5;
```

```
SELECT avg(val)
FROM avg_tests;
```

```
SELECT avg(DISTINCT val)
FROM avg_tests;
```

# SQLite Aggregate Functions - Max/MIN/SUM

The SQLite MAX/MIN/SUM function is an aggregate function that returns the maximum/minimum/sum value of all values in a group.

```
SELECT MAX(val)
FROM avg_tests;
```

```
SELECT MIN(val)
FROM avg_tests;
```

```
CREATE TABLE avg_tests (val);1
INSERT INTO avg_tests (val)
VALUES (1), (2), (10.1), (20.5),
('8'), ('B'), (NULL), (x'0010'), (x'0011');
```

```
SELECT SUM(val)
FROM avg_tests;
```

```
SELECT rowid,
       val
FROM avg_tests;
```

# Lab 2.1.1: SQL

- Purpose:
  - Create account in **Mode Analytics**
  - To discover the basic features of SQL
- Tools and Resources:
  - Mode Analytics (account required)
- Materials:
  - Mode Analytics interactive SQL tutorial  
<https://community.modeanalytics.com/sql/tutorial/introduction-to-sql/>





# Database Scripting

- shell scripts / commands
  - for quick / convenient execution of routine tasks
    - **populating / updating a database** from a file
    - **backing up or restoring** a database
    - **dumping** a database object to a file
    - **executing queries** to deliver rowsets for subsequent analysis
    - **moving data** between databases, data lakes, etc.
- examples

```
psql -U postgres -d database_name -c "SELECT c_defaults FROM user_info  
WHERE c_uid = 'testuser'"
```

```
mysql -h "server-name" -u "root" "-pXXXXXXXXX" "database-name" < "filename.sql"
```

# Database Administration

- granting **permissions**
  - users, roles
  - tables, views
- performing **backups & recoveries**
- creating & scheduling **jobs**
- Nb. It is not uncommon for a query to join tables from different databases
  - in SQL Server, the query must be created by the *single* owner of *both* databases
  - a database owner should be a virtual user -- not a particular person's login

# Discussion

- Why is the RDBMS still in use?
- QUESTIONS?

# Advanced SQL

- Aggregation functions
- Grouping
- Window functions

# SQL Aggregation functions

- aggregate many rows into a single resultant row
- common aggregation functions:
  - COUNT counts all rows meeting criterion
  - SUM sums selected field(s) in all rows meeting criterion
  - AVG averages ...
  - MIN, MAX computes minimum/maximum of ...
- SQL engine may support user-defined aggregation functions
- example:
 

```
SELECT MIN(sale_date) AS firstdt, MAX(sale_date) AS lastdt,  
SUM(sales) AS netsales FROM sales
```

# Aggregating by Groups in SQL

- grouping allows aggregation functions to be applied to subsets based on row-level criteria
- adds GROUP BY clause to select statement
- example:

```
SELECT agent_name, SUM(sales) AS netsales FROM sales
GROUP BY agent_name
```

```
SELECT COUNT(CustomerID), Country
```

```
FROM Customers
```

```
GROUP BY Country
```

```
ORDER BY COUNT(CustomerID) DESC;
```

# SQL Window Functions

- operates on a *set* of rows and return a value for *each* row
  - like aggregation, but performed in relation to current row
- example: running total

```
SELECT duration, SUM(duration)  
OVER (ORDER BY start_time) AS running_total  
FROM bikeshare
```

- OVER clause designates window
- ORDER BY sets sequence of rowset (oldest to newest values of start\_time)
- each row shows current `duration` and sum of all previous values of `duration`

# Lab 2.1.2: Advanced SQL

- Purpose:
  - To discover the more powerful features of SQL
- Tools and Resources:
  - Mode Analytics (account required)
  - Mode Analytics interactive SQL tutorial  
<https://community.modeanalytics.com/sql/tutorial/introduction-to-sql/>





# SQL in Python

- Python/SQL integration paradigms
- Python with embedded SQL db (SQLite)
- Python with external SQL db (pyodbc)

# Python / SQL Integration Paradigms

- **Embedded**

- an SQL RDBMS is emulated by a library designed to work with Python
- databases are not normally accessible outside of Python
- ideal for a self-contained Python app with its own db

- **External**

- a stand-alone SQL RDBMS resides on a database server that other applications can connect to
- the db is made accessible to Python via an ODBC driver specific to the RDBMS in use
  - SQL Server, MySQL, DB2, etc.
- used when the Python app is just a client of a more general-purpose db

# Python with Embedded SQL Database

## Example: SQLite

```
import sqlite3
connection = sqlite3.connect("company.db")
```

```
sql_command = """
CREATE TABLE employee (
staff_number INTEGER PRIMARY KEY,
fname VARCHAR(20),
lname VARCHAR(30),
date_joined DATE);"""
```

# Python with External SQL Database

## Example: pyodbc

```
import pyodbc
cnxn = pyodbc.connect("DSN=MSSQL-PYTHON")
cursor = cnxn.cursor()
cursor.tables()
rows = cursor.fetchall()
for row in rows:
    print row.table_name
```

# Python with pyodbc

- Python uses a cursor to iterate through a returned rowset

```
import pyodbc
import pandas.io.sql as psql
import pandas as pd

cxnstr = "Server=myServerAddress;Database=myDB;User Id=myUsername;Password=myPass;"
cxn = pyodbc.connect(cxnstr)
cursor = cxn.cursor()
cursor.execute("""SELECT ID, FirstName, LastName FROM mytable""")
rows = cursor.fetchall()
objects_list = []
for row in rows:
    d = collections.OrderedDict()
    d['UserID'] = row.ID
    d['FirstName'] = row.FirstName
    d['LastName'] = row.LastName
cxn.close()
```

# HOMEWORK

## 1. Install

- MongoDB Community Server <https://www.mongodb.com/download-center#community>
- Neo4j Community Server <https://neo4j.com/download-center/#releases>

## 2. Install Python packages:

- pymongo (conda)
- neo4j-driver (pip)

# Lab 2.1.3: SQL in Python

- Purpose:
  - To investigate SQL implementation in Python
- Tools and Resources:
  - Sqlit: Install on your labtop  
<https://www.sqlite.org/download.html>
  - Jupyter Notebooks
  - Python package sqlite3
- Materials:
  - 'Lab 2.1.3 – Databases.ipynb'
- Data:
  - 'housing-data.csv'



# Scalable SQL

- massively parallel processing relational database systems (**MPP** RDBMS)
  - (expensive) architecture supports scalability and high performance



- NoSQL databases *with SQL*



# Where Does Big Data Come From?

It's all happening online – could record every:

- » Click
- » Ad impression
- » Billing event
- » Fast Forward, pause,...
- » Server request
- » Transaction
- » Network message
- » Fault
- » ...




# Where Does Big Data Come From?

User Generated Content (Web & Mobile)

- » Facebook
- » Instagram
- » Yelp
- » TripAdvisor
- » Twitter
- » YouTube
- » ...

# Where Does Big Data Come From?

## Health and Scientific Computing




# Where Does Big Data Come From?

## Graph Data

Lots of interesting data has a graph structure:

- Social networks
- Telecommunication Networks
- Computer Networks
- Road networks
- Collaborations/Relationships
- ...

Some of these graphs can get quite large  
(e.g., Facebook user graph)



# Where Does Big Data Come From?

## Log Files – Apache Web Server Log

```
uplherc.upl.com - - [01/Aug/1995:00:00:07 -0400] "GET / HTTP/1.0" 304 0
uplherc.upl.com - - [01/Aug/1995:00:00:08 -0400] "GET /images/ksclogo-medium.gif HTTP/1.0" 304 0
uplherc.upl.com - - [01/Aug/1995:00:00:08 -0400] "GET /images/MOSAIC-logosmall.gif HTTP/1.0" 304 0
uplherc.upl.com - - [01/Aug/1995:00:00:08 -0400] "GET /images/USA-logosmall.gif HTTP/1.0" 304 0
ix-esc-ca2-07.ix.netcom.com - - [01/Aug/1995:00:00:09 -0400] "GET /images/launch-logo.gif HTTP/1.0" 200 1713
uplherc.upl.com - - [01/Aug/1995:00:00:10 -0400] "GET /images/WORLD-logosmall.gif HTTP/1.0" 304 0
slppp6.intermind.net - - [01/Aug/1995:00:00:10 -0400] "GET /history/skylab/skylab.html HTTP/1.0" 200 1687
piweba4y.prodigy.com - - [01/Aug/1995:00:00:10 -0400] "GET /images/launchmedium.gif HTTP/1.0" 200 11853
tampico.usc.edu - - [14/Aug/1995:22:57:13 -0400] "GET /welcome.html HTTP/1.0" 200 1
```

# Where Does Big Data Come From?

Internet of Things: RFID tags

California FasTrak Electronic Toll Collection transponder

Used to pay tolls

Collected data also  
used for traffic reporting  
» <http://www.511.org/>



# The Big Data Problem

Data growing faster than computation speeds

Growing data sources

» Web, mobile, scientific, ...

Storage getting cheaper

» Size doubling every 18 months

But, stalling CPU speeds and storage bottlenecks

Big Data Examples

Facebook's daily logs: 60 TB

1,000 genomes project: 200 TB

Google web index: 10+ PB


Cost of 1 TB of disk: ~\$35

Time to read 1 TB from disk: 3 hours  
(100 MB/s)




# The Big Data Problem

One machine can not process or even store all the data!  
Solution is to distribute data over cluster of machines



Lots of hard drives



... and CPUs

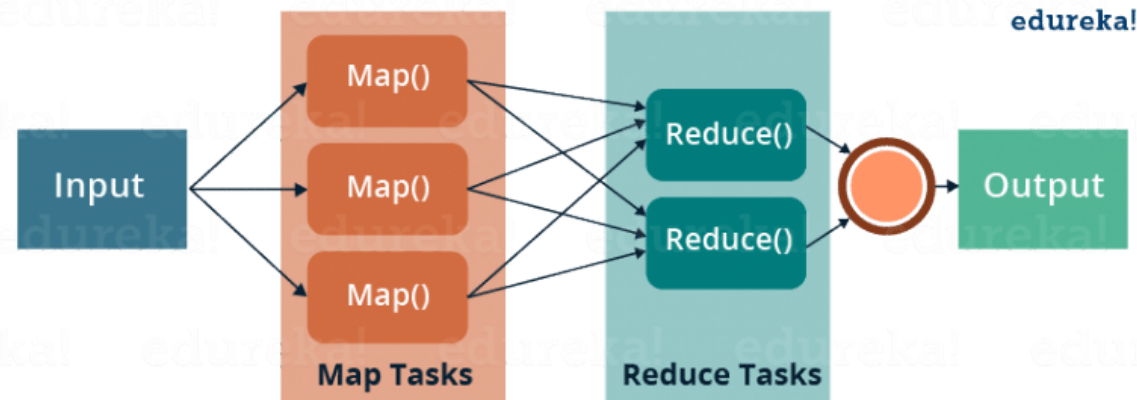


... and memory!

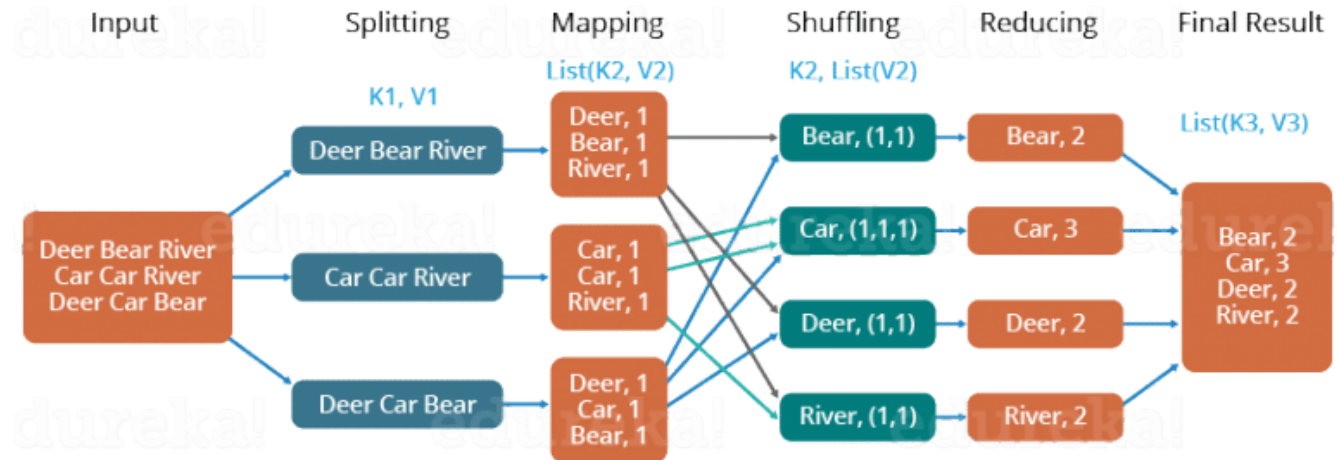


# MapReduce

MapReduce is a programming framework that allows us to perform distributed and parallel processing on large data sets in a distributed environment.



The Overall MapReduce Word Count Process



# Motivation

Most current cluster programming models are based on *acyclic data flow* from stable storage to stable storage

**Benefits of data flow:** runtime can decide where to run tasks and can automatically recover from failures

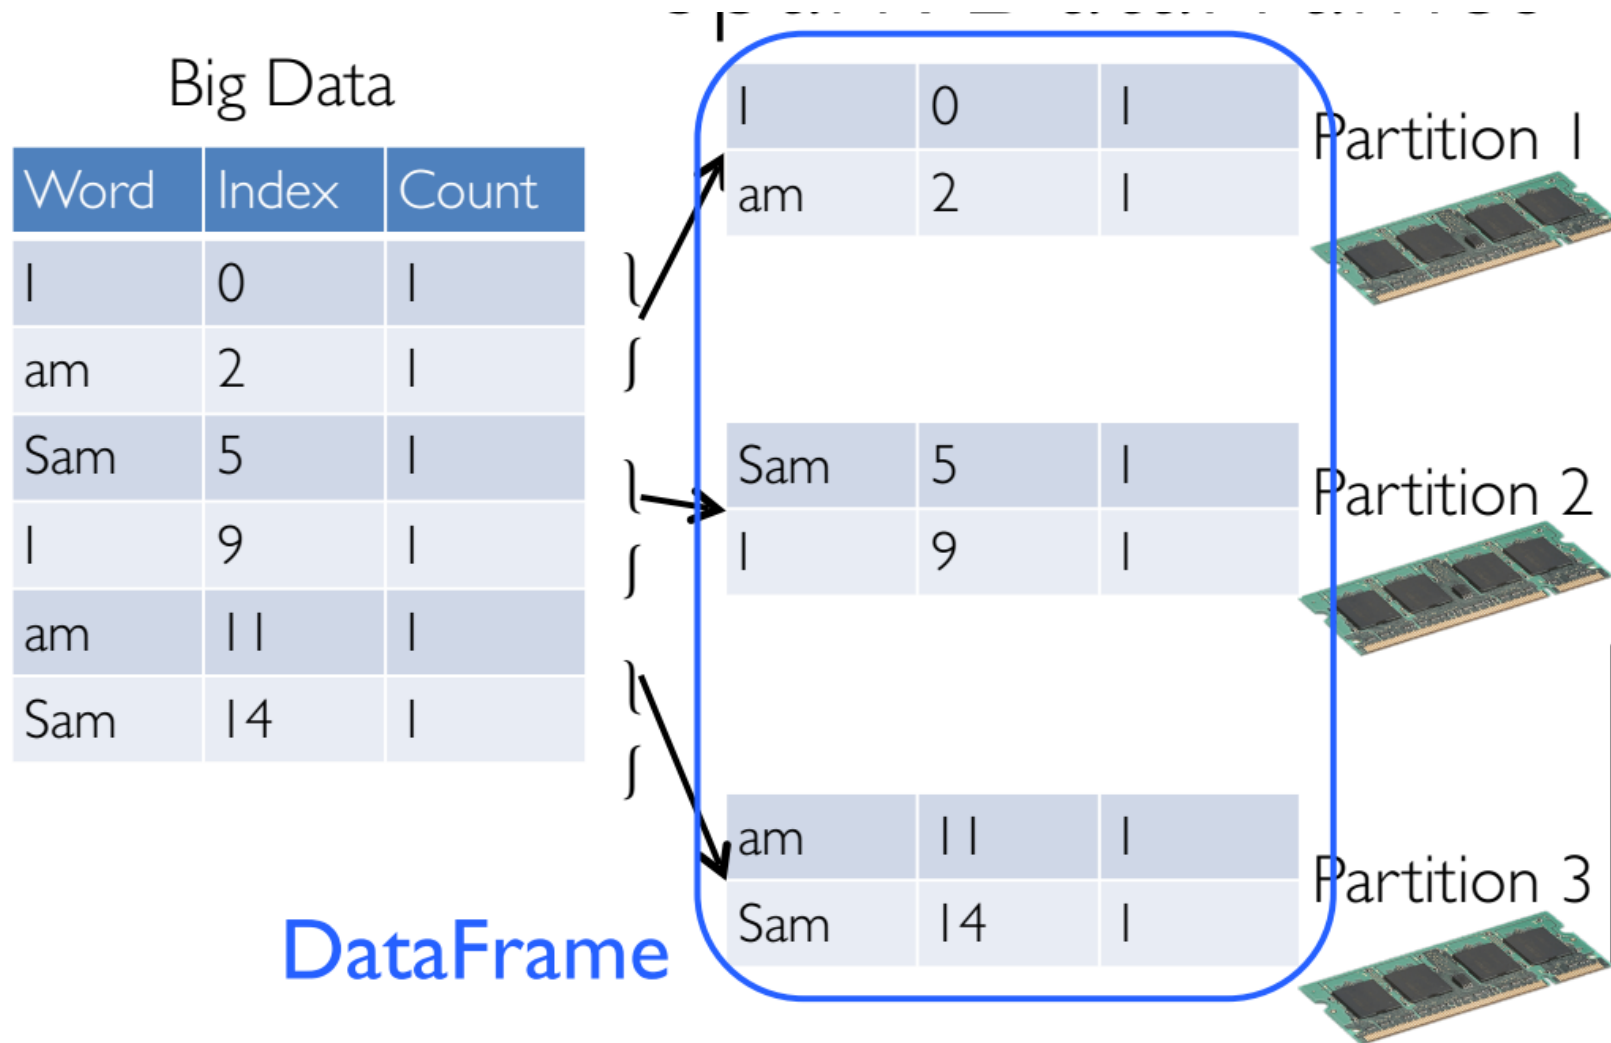
Acyclic data flow is inefficient for applications that repeatedly reuse a *working set* of data:

**Iterative** algorithms (machine learning, graphs)

**Interactive** data mining tools (R, Excel, Python)

With current frameworks, apps reload data from stable storage on each query

# Spark DataFrames



# Spark - Fast, Interactive, Language-Integrated Cluster Computing

**Solution:** Resilient Distributed Datasets (RDDs)

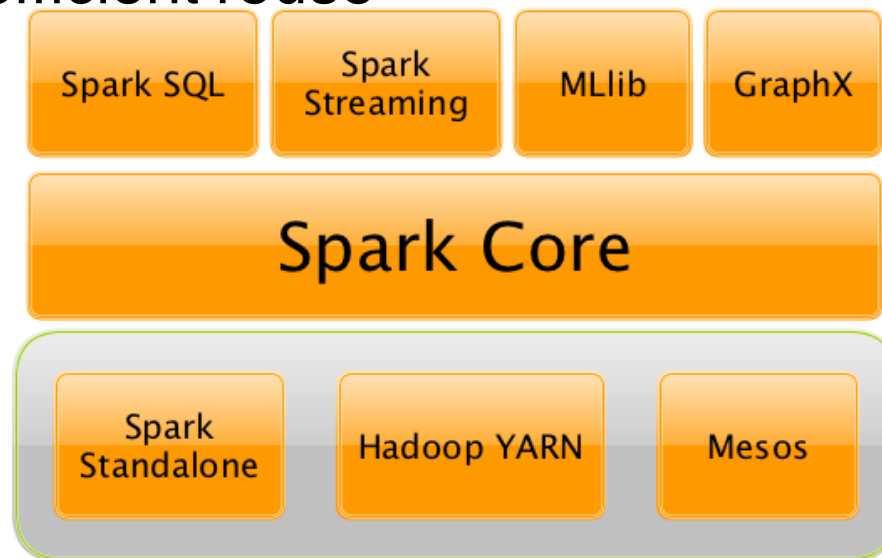
Spark is often called cluster computing engine or simply execution engine. You could also describe Spark as a distributed, data processing engine for batch and streaming modes featuring SQL queries, graph processing, and machine learning.

Allow apps to keep working sets in memory for efficient reuse

Retain the attractive properties of MapReduce

- Fault tolerance, data locality, scalability

Support a wide range of applications

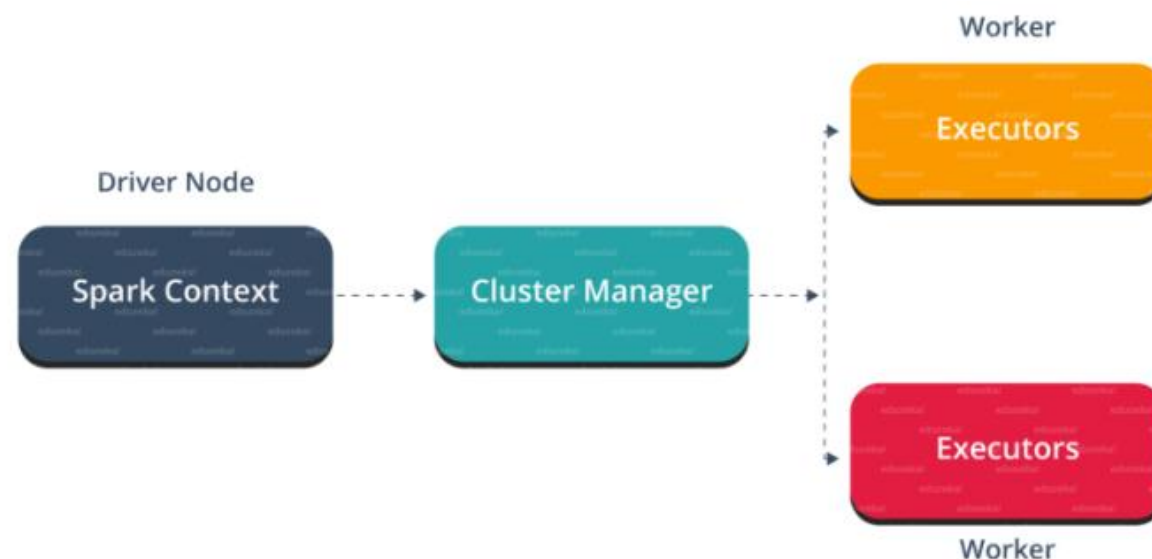


# Spark Architecture Overview

## Spark Architecture Overview

Apache Spark has a well-defined layered architecture where all the spark components and layers are loosely coupled. This architecture is further integrated with various extensions and libraries. Apache Spark Architecture is based on two main abstractions:

- Resilient Distributed Dataset (RDD)
- Directed Acyclic Graph (DAG)




# Spark - Fast, Interactive, Language-Integrated Cluster Computing

In contrast to Hadoop's **two-stage disk-based MapReduce computation engine**, Spark's **multi-stage (mostly) in-memory** computing engine allows for **running most computations in memory**, and hence most of the time provides better performance for certain applications, e.g. iterative algorithms or interactive data mining (read Spark officially sets a new record in large-scale sorting).

Spark aims at speed, ease of use, extensibility and interactive analytics.

Spark is a distributed platform for executing complex multi-stage applications, like machine learning algorithms, and interactive ad hoc queries. Spark provides an efficient abstraction for in-memory cluster computing called **Resilient Distributed Dataset**.

# Logistic Regression Performance




# Spark Applications

- In-memory data mining on Hive data (Conviva)
- Predictive analytics (Quantifind)
- City traffic prediction (Mobile Millennium)
- Twitter spam classification (Monarch)
- Collaborative filtering via matrix factorization



# Conviva GeoReport



Aggregations on many keys w/ same WHERE clause

40 × gain comes from:

- » Not re-reading unused columns or filtered records
- » Avoiding repeated decompression
- » In-memory storage of deserialized objects

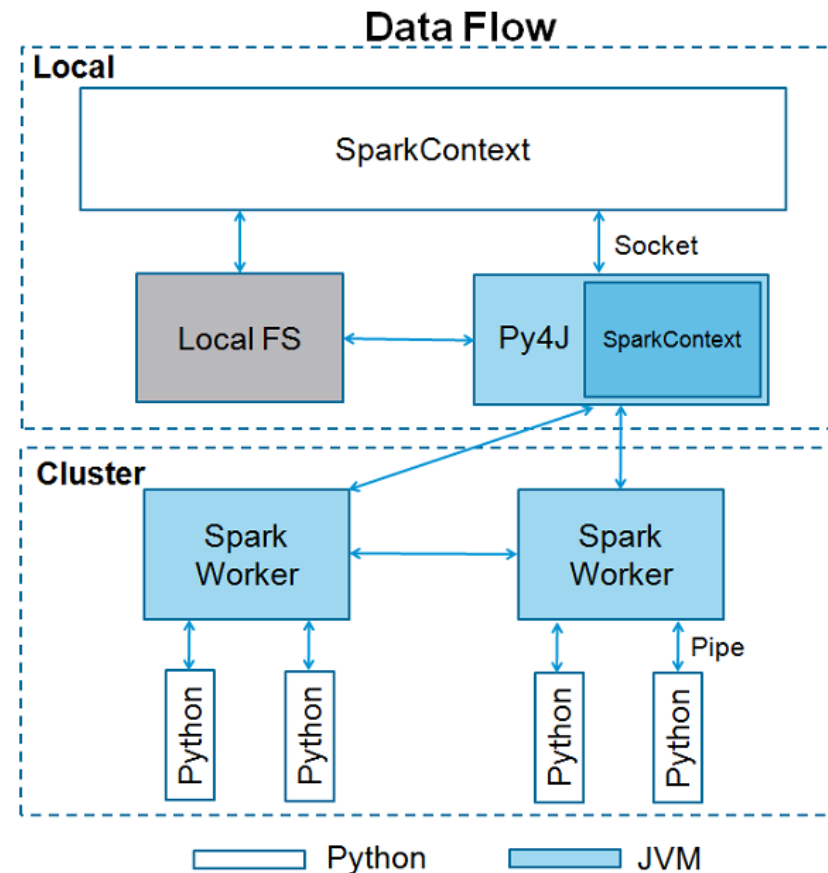
# Python Spark (pySpark)

We are using the Python programming interface to Spark (pySpark)

pySpark provides an easy-to-use programming abstraction and parallel runtime:

» “Here’s an operation, run it on all of the data”

DataFrames are the key concept



# Spark and SQL Contexts

A Spark program first creates a SparkContext object

- » SparkContext tells Spark how and where to access a cluster
- » pySpark shell, Databricks CE automatically create SparkContext
- » iPython and programs must create a new SparkContext

The program next creates a sqlContext object

Use sqlContext to create DataFrames

# SparkContext Example - Python Program

```
from pyspark import SparkContext
logFile = "file:///home/hadoop/spark-2.1.0-bin-hadoop2.7/README.md"
sc = SparkContext("local", "first app")
logData = sc.textFile(logFile).cache()
numAs = logData.filter(lambda s: 'a' in s).count()
numBs = logData.filter(lambda s: 'b' in s).count()
print "Lines with a: %i, lines with b: %i" % (numAs, numBs)
```

# SparkContext Example - Python Program

Basically, this program just counts the number of lines in 'a' and 'b' in a text file. However, we need to replace \$YOUR\_SPARK\_HOME with the Spark's installation location.

```
from pyspark import SparkContext
```

```
logFile = "$YOUR_SPARK_HOME/README.md" # Should be some file on your system
```

```
sc = SparkContext("local", "Simple App1")
```

```
logData = sc.textFile(logFile).cache()
```

```
numAs = logData.filter(lambda s: 'a' in s).count()
```

```
numBs = logData.filter(lambda s: 'b' in s).count()
```

```
print "Lines with a: %i, lines with b: %i" % (numAs, numBs)
```

# Continue PySpark on github

<https://github.com/andfanilo/pyspark-tutorial>

<https://github.com/PacktPublishing/Learning-PySpark>

# NoSQL Databases

- Scaling Up
- What (and why) are NoSQL databases
- Characteristics of NoSQL databases
- NoSQL database types

# Scaling Up

- Issues with scaling up when the dataset is just too big
- RDBMS were not designed to be distributed
- Began to look at multi-node database solutions
- Known as ‘scaling out’ or ‘horizontal scaling’
- Different approaches include:
  - Master-slave
  - Sharding



# Scaling RDBMS - Sharding

- Partition or sharding
  - Scales well for both reads and writes
  - Not transparent, application needs to be partition-aware
  - Can no longer have relationships/joins across partitions
  - Loss of referential integrity across shards

# Other ways to scale RDBMS

- Multi-Master replication
- INSERT only, not UPDATES/DELETES
- No JOINS, thereby reducing query time
  - This involves de-normalizing data
- In-memory databases

# What is NoSQL?

- Stands for **Not Only SQL**
- Class of non-relational data storage systems
- Usually do not require a fixed table schema nor do they use the concept of joins
- All NoSQL offerings relax one or more of the ACID properties (will talk about the CAP theorem)

# Why NoSQL?

- For data storage, an RDBMS cannot be the be-all/end-all
- Just as there are different programming languages, need to have other data storage tools in the toolbox
- A NoSQL solution is more acceptable to a client now than even a year ago
  - Think about proposing a Ruby/Rails or Groovy/Grails solution now versus a couple of years ago

# How did we get here?

- Explosion of social media sites (Facebook, Twitter) with large data needs
- Rise of cloud-based solutions such as Amazon S3 (simple storage solution)
- Just as moving to dynamically-typed languages (Ruby/Groovy), a shift to dynamically-typed data with frequent schema changes
- Open-source community

# Dynamo and BigTable

- Three major papers were the seeds of the NoSQL movement
  - BigTable (Google)
  - Dynamo (Amazon)
    - Gossip protocol (discovery and error detection)
    - Distributed key-value data store
    - Eventual consistency
  - CAP Theorem (discuss in a sec ..)

# The Perfect Storm

- Large datasets, acceptance of alternatives, and dynamically-typed data has come together in a perfect storm
- Not a backlash/rebellion against RDBMS
- SQL is a rich query language that cannot be rivaled by the current list of NoSQL offerings

# CAP Theorem

- Three properties of a system: consistency, availability and partitions
- You can have at most two of these three properties for any shared-data system
- To scale out, you have to partition. That leaves either consistency or availability to choose from
- In almost all cases, you would choose availability over consistency



# What kinds of NoSQL

- NoSQL solutions fall into two major areas:
  - Key/Value or 'the big hash table'.
    - Amazon S3 (Dynamo)
    - Voldemort
    - Scalaris
  - Schema-less which comes in multiple flavors, column-based, document-based or graph-based.
    - Cassandra (column-based)
    - CouchDB (document-based)
    - Neo4J (graph-based)
    - HBase (column-based)

# Key/Value

## *Pros:*

- very fast
- very scalable
- simple model
- able to distribute horizontally

## *Cons:*

- many data structures (objects) can't be easily modeled as key value pairs

# Schema-Less

## *Pros:*

- Schema-less data model is richer than key/value pairs
- eventual consistency
- many are distributed
- still provide excellent performance and scalability

## *Cons:*

- typically no ACID transactions or joins

# Characteristics of NoSQL Databases

- Cheap, easy to implement (open source)
- Data are replicated to multiple nodes (therefore identical and fault-tolerant) and can be partitioned
  - Down nodes easily replaced
  - No single point of failure
- Easy to distribute
- Don't require a schema
- Can scale up and down
- Relax the data consistency requirement (



# Characteristics of NoSQL Databases – cont'd

## *Disadvantages*

- many don't support true ACID transactions
  - application code is obliged to try to manage concurrency issues
- easy to modify schemas
  - application developers need to collaborate closely to ensure schema development is under control
- much slower than RDBMS
- lack powerful management & development features of RDBMS

# Use Cases

- simple data requirements
- application-specific data
- start-ups



# What are we giving up?

- joins
- group by
- order by
- ACID transactions
- SQL as a sometimes frustrating but still powerful query language
- easy integration with other applications that support SQL

# Cassandra

- Originally developed at Facebook
- Follows the BigTable data model: column-oriented
- Uses the Dynamo Eventual Consistency model
- Written in Java
- Open-sourced and exists within the Apache family
- Uses Apache Thrift as it's API



# Thrift

- Created at Facebook along with Cassandra
- Is a cross-language, service-generation framework
- Binary Protocol (like Google Protocol Buffers)
- Compiles to: C++, Java, PHP, Ruby, Erlang, Perl, ...

# Searching

- Relational
  - `SELECT `column` FROM `database`,`table` WHERE `id` = key;`
  - `SELECT product_name FROM rockets WHERE id = 123;`
- Cassandra (standard)
  - `keyspace.getSlice(key, "column_family", "column")`
  - `keyspace.getSlice(123, new ColumnParent("rockets"), getSlicePredicate());`

# Typical NoSQL API

- Basic API access:
  - `get(key)` -- Extract the value given a key
  - `put(key, value)` -- Create or update the value given its key
  - `delete(key)` -- Remove the key and its associated value
  - `execute(key, operation, parameters)` -- Invoke an operation to the value (given its key) which is a special data structure (e.g. List, Set, Map .... etc).

# NoSQL Databases Types

- wide column store
  - > Hadoop / **HBase**, MapR, BigTable, Hortonworks, Cloudera, Cassandra, Informix
- document store
  - > MongoDB, CouchDB, Azure DocumentDB
- key value / tuple store
  - > DynamoDB, Azure Table Storage, Oracle NoSQL
- graph databases
  - > Neo4j
- multi-model databases
  - > ArangoDB, OrientDB
- object databases
  - > Versant, Objectivity VelocityDB



## NoSQL Databases – cont'd

- Document databases
  - MongoDB

# Document Databases

- semi-structured data
- records do not all need to have the same fields

## XML:

a record is a block of XML tags and values

```
<contact>
  <firstname>Bob</firstname>
  <address>5 Oak St.</lastname>
  <hobby>saling</hobby>
</contact>
```

## JSON:

a record is a list of key:value pairs

```
{
  "FirstName": "Bob",
  "Address": "5 Oak St.",
  "Hobby": "sailing"
}
```

# MongoDB

- an **open-source** document database
  - no charge for *Community Server* version
- high performance
  - **embedded data models** reduce I/O activity
  - indexes
    - can include keys from embedded documents, arrays
- high availability
  - replication facility
    - automatic fail-over
    - data redundancy
- automatic scalability (horizontal)



# MongoDB

- CRUD ✓
  - create, read, update, delete
- ACID ?
  - traditionally:
    - document databases are only ACID-compliant only at document level
    - no *transactions* for containing multiple I/O operations
    - application code is obliged to emulate transactions, if required
    - latency results in an **eventual consistency** model
  - MongoDB 4.0
    - introduced transactions



# MongoDB: High-Level Objects

## *Document:*

- a set of field:value pairs
  - values can be hierarchical
- analogous to RDB row

examples:

```
{ name: "sue", age: 26, status: "A", groups: [ "news", "sports" ] }
```

```
{ name: "fred", status: "A", groups: [ "sports", "hobbies", "cars" ] }
```

```
{ name: { first: "fred", last: "bloggs" }, status: "A", groups: [ "sports", "hobbies", "cars" ] }
```

## *Collection:*

- a logical group of documents
- analogous to RDB table

# MongoDB with Python

## example:

```
import pymongo
connection =
pymongo.MongoClient("mongodb://localhost")
db = connection.school
students = db.students
cursor = students.find()
```

```
# find minimum homework score...
for doc in cursor:
    scores = doc["scores"]
    minhs = 101
    for entry in scores:
        if entry["type"] == "homework":
            if entry["score"] < minhs:
                minhs = entry["score"]
```

- create mongod client
- connect to database 'school'
- create alias for table 'students'
- fetch all rows into cursor
- loop through docs in cursor
- get value (doc) associated with key 'scores'
- initialise min to impossibly large value (> 100%)
- loop through 'scores' docs, looking for 'homework' keys
- test each corresponding score to find new minimum

## Lab 2.1.4: Python with MongoDB (Optional homework)

- Purpose:
  - To develop skills in NoSQL database programming with MongoDB
- Materials:
  - 'Lab 2.1.4.ipynb'





## NoSQL Databases – cont'd

- Graph Databases
  - Neo4j

# Graph Databases

- model members and relationships as a network
- high-level objects:
  - nodes
    - entities (e.g. people, accounts, organisations)
  - edges
    - connections between nodes
  - properties
    - node: differentiates types of nodes
      - roles, classifications, etc.
    - edge: describes the relationship
      - 2-way, 1-way, directionless
      - friend, follower, commenter, etc.

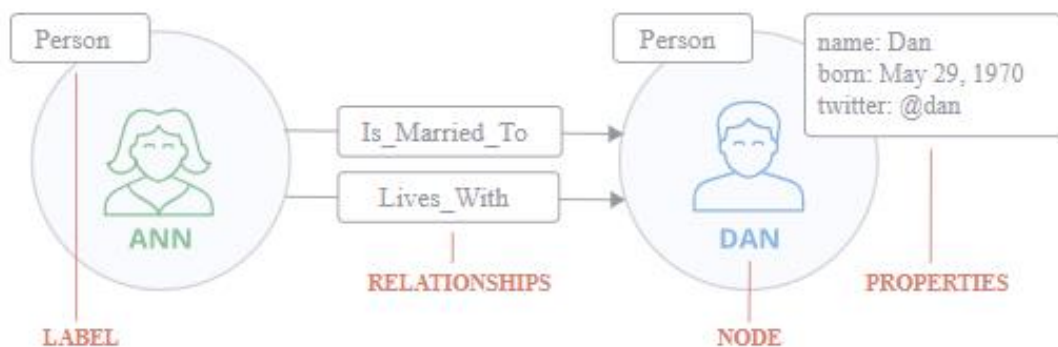
# Neo4j

- open source
  - no charge for *Community Server* edition
- ACID-compliant, transactional database with native graph storage & processing
- online backup
- high availability
- most popular graph database



# Neo4j: Basics

## The Labeled Property Graph Model




### Nodes

- Nodes are the main data elements
- Nodes are connected to other nodes via **relationships**
- Nodes can have one or more **properties** (i.e., attributes stored as key/value pairs)
- Nodes have one or more **labels** that describes its role in the graph

### Relationships

- Relationships connect two nodes
- Relationships are directional
- **Nodes** can have multiple, even recursive relationships
- Relationships can have one or more **properties** (i.e., attributes stored as key/value pairs)

# Cypher Query Language



- example: find friends of friends of John

```
MATCH (john {name: 'John'})-[:friend]->()  
()-[:friend]->(fof)  
RETURN john.name, fof.name
```

output:

+-----+	
john.name	fof.name
+-----+	
"John"	"Maria"
"John"	"Steve"
+-----+	



## Lab 2.1.5: Neo4j and Python (Optional homework)

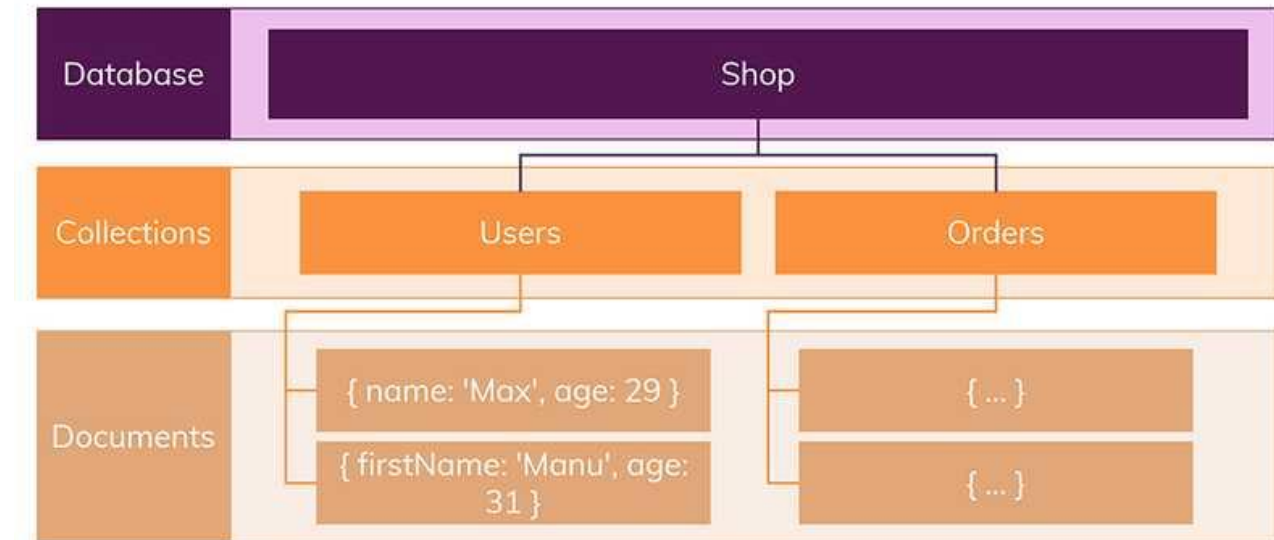
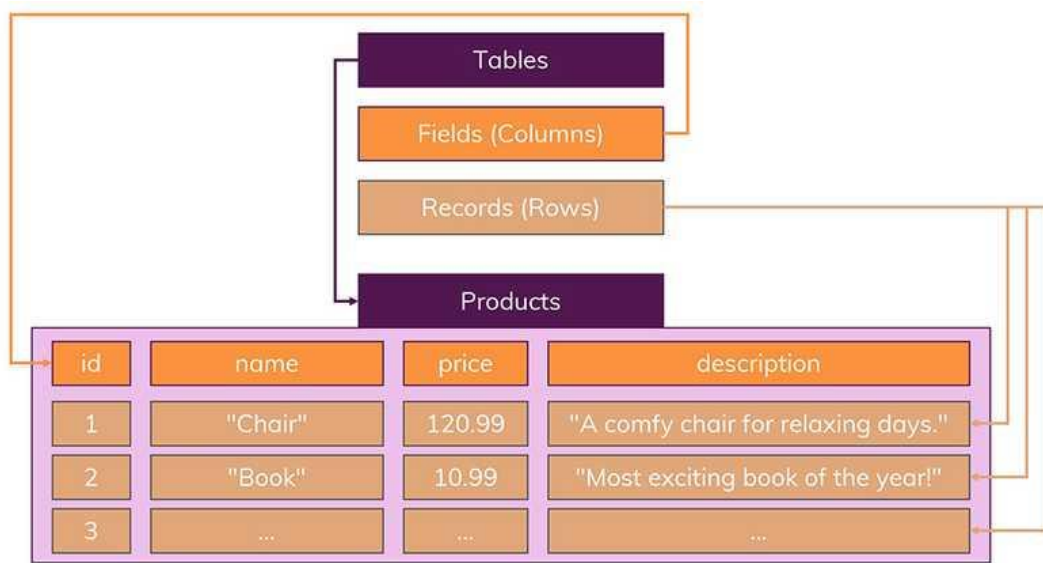
- Purpose:
  - To develop familiarity with graph database programming (Neo4j) using:
    - the Neo4j GUI
    - a Python library for Neo4j
- Resources:
  - Neo4j built-in tutorials
  - Cypher cheatsheet
    - <https://neo4j.com/docs/cypher-refcard/3.2/>
- Materials:
  - 'Lab 2.1.5.ipynb'



# Discussion: SQL vs NoSQL

SQL	NoSQL
Traditional rows and columns governed data model	No predefined data structure database at mercy of developers
Strict structure (incl. primary keys) schema changes difficult, risky	Ideal for unstructured data schema can change with application requirements
Entire column for each feature	Cheaper hardware
Industry standard	Supports design flexibility & growth popular among startups
ACID	Application code must manage transactions

# Discussion: SQL vs NoSQL



# Discussion: SQL vs NoSQL

Users		
id	email	name
1	max@test.com	Maximilian Schwarzmüller
2	manu@test.com	Manuel Lorenz
3	...	...

Products			
id	title	price	description
1	Chair	120.99	A comfy chair
2	Book	10.99	Exciting book!
3	...	...	...

Orders		
id	user_id	product_id
1	2	1
2	1	1
3	2	2

Orders		Duplicate Data
{ id: 'ddjfa31', user: { id: 1, email: 'max@test.com' }, product: { id: 2, price: 10.99 } }		
{ id: 'lddao1', user: { id: 2, email: 'manu@test.com' }, product: { id: 1, price: 120.99 } }		
{ id: 'nbax12', product: { id: 2, price: 10.99 } }		
{ ... }		


Users	
{ id: 1, name: 'Max', email: 'max@test.com' }	
{ id: 2, name: 'Manu', email: 'manu@test.com' }	
{ ... }	

Products	
{ id: 1, title: 'Chair', price: 120.99 }	
{ id: 2, name: 'Book', price: 10.99 }	
{ ... }	


# Discussion: SQL vs NoSQL

Horizontal Scaling



Add More Servers (and merge Data into one Database)

Vertical Scaling



Improve Server Capacity / Hardware

# Discussion: NoSQL with SQL ?!

- Why has SQL infiltrated the NoSQL paradigm?

# Questions?

# Appendices




# Relational Databases – Normalisation

## Codd's 1st-normal form

- the domain of each attribute contains only atomic (indivisible) values
- the value of each attribute contains only a single value from that domain

Customer			
Customer ID	First Name	Surname	Telephone Number
123	Pooja	Singh	555-861-2025, 192-122-1111
456	San	Zhang	(555) 403-1659 Ext. 53; 182-929-2929
789	John	Doe	555-808-9633



Customer			
Customer ID	First Name	Surname	Telephone Number
123	Pooja	Singh	555-861-2025
123	Pooja	Singh	192-122-1111
456	San	Zhang	182-929-2929
456	San	Zhang	(555) 403-1659 Ext. 53
789	John	Doe	555-808-9633




# Relational Databases – – Normalisation - cont'd

## Codd's 2nd-normal form

- in 1st-normal form
- no non-prime attribute is dependent on any proper subset of any candidate key of the relation  
(an attribute that is not a part of any candidate key of the relation)

**Electric Toothbrush Models**

<u>Manufacturer</u>	<u>Model</u>	Model Full Name	Manufacturer Country
Forte	X-Prime	Forte X-Prime	Italy
Forte	Ultraclean	Forte Ultraclean	Italy
Dent-o-Fresh	EZbrush	Dent-o-Fresh EZbrush	USA
Kobayashi	ST-60	Kobayashi ST-60	Japan
Hoch	Toothmaster	Hoch Toothmaster	Germany
Hoch	X-Prime	Hoch X-Prime	Germany



**Electric Toothbrush Manufacturers**

<u>Manufacturer</u>	Manufacturer Country
Forte	Italy
Dent-o-Fresh	USA
Kobayashi	Japan
Hoch	Germany

**Electric Toothbrush Models**


<u>Manufacturer</u>	<u>Model</u>	Model Full Name
Forte	X-Prime	Forte X-Prime
Forte	Ultraclean	Forte Ultraclean
Dent-o-Fresh	EZbrush	Dent-o-Fresh EZbrush
Kobayashi	ST-60	Kobayashi ST-60
Hoch	Toothmaster	Hoch Toothmaster
Hoch	X-Prime	Hoch X-Prime

# Relational Databases – – Normalisation - cont'd

## Codd's 3rd-normal form

- in 2nd-normal form
- every non-prime attribute (of a table) is non-transitively dependent on every key (of a table)

Tournament Winners			
<u>Tournament</u>	<u>Year</u>	Winner	Winner Date of Birth
Indiana Invitational	1998	Al Fredrickson	21 July 1975
Cleveland Open	1999	Bob Albertson	28 September 1968
Des Moines Masters	1999	Al Fredrickson	21 July 1975
Indiana Invitational	1999	Chip Masterson	14 March 1977



Tournament Winners			Winner Dates of Birth	
<u>Tournament</u>	<u>Year</u>	Winner	<u>Winner</u>	Date of Birth
Indiana Invitational	1998	Al Fredrickson	Chip Masterson	14 March 1977
Cleveland Open	1999	Bob Albertson	Al Fredrickson	21 July 1975
Des Moines Masters	1999	Al Fredrickson	Bob Albertson	28 September 1968
Indiana Invitational	1999	Chip Masterson		



# Which RDBMS Object to Use When

- queries
  - **ad hoc** queries
  - **stored** or **generated** in application code
- views
  - stored (**reusable**) queries
  - can incorporate **joins** with other views
  - **preferable** to queries in application code
- stored procedures
  - **more powerful than views** (can query and/or modify data)
  - preferred for delivering data to applications (**security, control, maintainability**)
- reports
  - **formatted output** containers based on tables, views
  - text & graphics
  - usually provided via a separate application (designed for but existing outside the RDBMS)
  - may have built-in subscription service

# End of presentation